

Why is Hadoop's block size set to 64MB by default, when most filesystems have block sizes of 16KB or less?

- ☒ ] If Hadoop's block size was set to 16KB, there would be a huge number of blocks throughout the cluster, which causes the NameNode to manage an enormous amount of metadata
- ☒ ] Since we need a Mapper for each block that we want to process, there would be a lot of Mappers, each processing a piece bit of data, which isn't efficient
- ☐ ] Because you can only store one block per node on the cluster
- ☐ ] Because a 16KB block is too small for multiple Mappers to process simultaneously